

## comp429-Project 2

---

### Group Members:

- M. Musab Küçük (mkucuk18-69910)
- Yavuzhan Akyüz (yakyuz18-68646)

GPU used for testing: TESLA V100

V1: works!

V2: works!

V3: works!

V4: works!

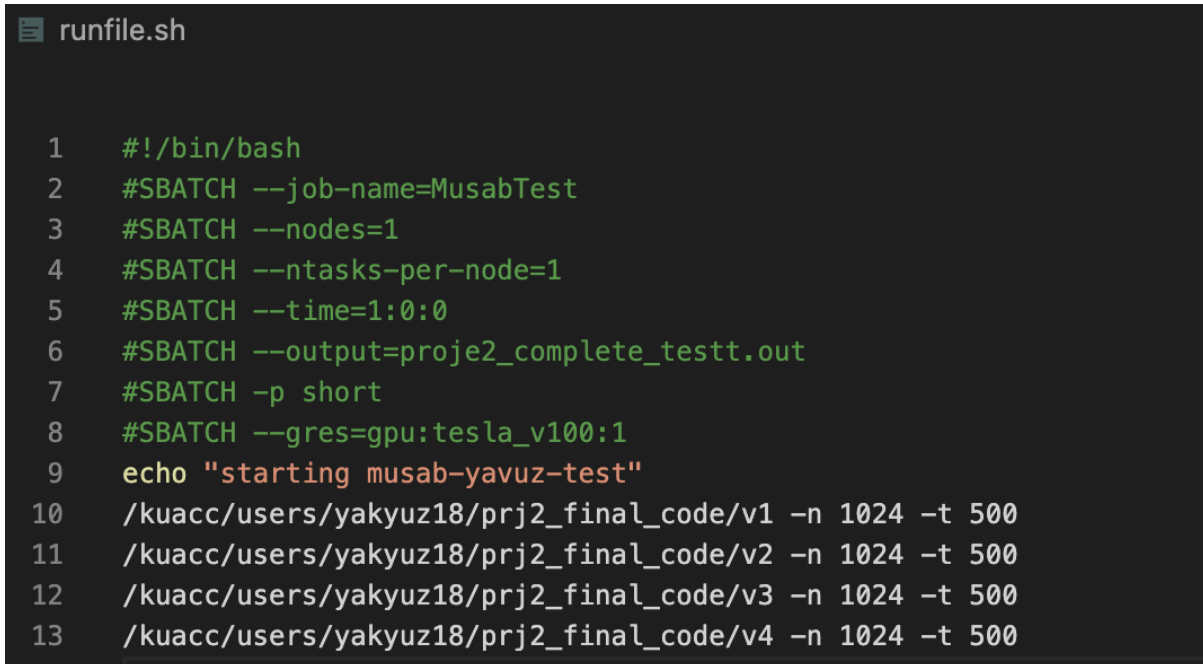
### How to run our code:

We have changed our Makefile, so just compile the code with "make" and then there will be executables named: serial, v1, v2, v3, v4 as seen on the Figure 1.

```
[[yakyuz18@login02 prj2_final_code]$ pwd
/scratch/users/yakyuz18/prj2_final_code
[[yakyuz18@login02 prj2_final_code]$ ls -l
Makefile
bin
cardiacsim.cpp
cardiacsim.o
cardiacsim_kernels.cu
cardiacsim_v1.o
cardiacsim_v2.o
cardiacsim_v3.o
cardiacsim_v4.o
proje2_complete_testt.out
runfile.sh
serial
v1
v1.cu
v2
v2.cu
v3
v3.cu
v4
v4.cu
[[yakyuz18@login02 prj2_final_code]$
```

Figure 1

-Now it's possible to use "srun" for each of these 1 by 1, or you can use the "sbatch runfile.sh", but please edit the path in "runfile.sh" as seen the figure 2



```
runfile.sh

1  #!/bin/bash
2  #SBATCH --job-name=MusabTest
3  #SBATCH --nodes=1
4  #SBATCH --ntasks-per-node=1
5  #SBATCH --time=1:0:0
6  #SBATCH --output=proje2_complete_testt.out
7  #SBATCH -p short
8  #SBATCH --gres=gpu:tesla_v100:1
9  echo "starting musab-yavuz-test"
10 /kuacc/users/yakyuz18/prj2_final_code/v1 -n 1024 -t 500
11 /kuacc/users/yakyuz18/prj2_final_code/v2 -n 1024 -t 500
12 /kuacc/users/yakyuz18/prj2_final_code/v3 -n 1024 -t 500
13 /kuacc/users/yakyuz18/prj2_final_code/v4 -n 1024 -t 500
```

Figure 2

## Implementation

---

### **Version 1:**

- We parallelized the serial cpp code in a new v1.cu file. We mostly used the same code, but made critical changes in our simulate function, and our Alloc2D function(we didn't use it to create arrays). Also we made sure to implement correct memory allocations and copies for device.

## Changes in Simulate:

- In order to parallelized serial program, we removed all the for loops in our simulate function, and rather we created three `__global__` kernel functions: `odeKernel`, `pdeKernel`, `ghostKernel`. Then for our every array ( `E` , `R`, `E_prev` ) we filled our new kernels accordingly. See Figure 3 for our implementation.

```

86  __global__ void ghostKernel(double *E_prev, const int n, const int m) {
87      int j = threadIdx.x + 1;
88
89      E_prev[j * (n+2)] = E_prev[j * (n+2) + 2];
90      E_prev[j * (n+2) + (n + 1)] = E_prev[j * (n + 2) + (n - 1)];
91
92      E_prev[j] = E_prev[2 * (n + 2) + j];
93      E_prev[(m + 1) * (n + 2) + j] = E_prev[(m - 1) * (n + 2) + j];
94  }
95
96  __global__ void odeKernel(double *E, double *R,
97      const int n, const int m, const double kk,
98      const double dt, const double a, const double epsilon,
99      const double M1, const double M2, const double b) {
100
101      int i = threadIdx.x + 1;
102      int j = blockIdx.x + 1;
103      int index = j * (n + 2) + i;
104
105      E[index] = E[index] - dt * (kk * E[index] * (E[index] - a) * (E[index] - 1) + E[index] * R[index]);
106      R[index] = R[index] + dt * (epsilon + M1 * R[index] / (E[index] + M2)) * (-R[index] - kk * E[index] * (E[index] - b - 1));
107  }
108
109  __global__ void pdeKernel(double *E, double *E_prev, const double alpha, const int n, const int m) {
110      int i = threadIdx.x + 1;
111      int j = blockIdx.x + 1;
112      int index = j * (n + 2) + i;
113
114      E[index] = E_prev[index] + alpha *
115          (E_prev[index + 1] + E_prev[index - 1] - 4 * E_prev[index] + E_prev[index + m + 2] +
116           E_prev[index - (m + 2)]);
117  }
118
119  void simulate(double *E, double *E_prev, double *R,
120      const double alpha, const int n, const int m, const double kk,
121      const double dt, const double a, const double epsilon,
122      const double M1, const double M2, const double b) {
123
124      ghostKernel<<<1, n>>>>(E_prev, n, m);
125      pdeKernel<<<m, n>>>>(E, E_prev, alpha, n, m);
126      odeKernel<<<m, n>>>>(E, R, n, m, kk, dt, a, epsilon, M1, M2, b);
127  }
128

```

Figure 3

**Changes in Memory Allocation:**

- So far in classes, we have dealt with 1D arrays for Memcpy to device, but in this project our main arrays were allocated in Host by 2D Alloc method, but we couldn't copy these array to our Device's memory, because we think in CUDA Memcpy must do linear memory allocation. Thus we converted our Host arrays to 1D arrays, so in all of our project we work with 2D -> 1D arrays. Because of our one cell padding, we created Device arrays with size (m+2, n+2). Then we used cudaMemcpy to copy them to Device , see Figure 4.

```
202  cudaMalloc((void **) &d_E, sizeof(double) * (m + 2) * (n + 2));
203  cudaMalloc((void **) &d_E_prev, sizeof(double) * (m + 2) * (n + 2));
204  cudaMalloc((void **) &d_R, sizeof(double) * (m + 2) * (n + 2));
205
206  cudaMemcpy(d_E, E, sizeof(double) * (m + 2) * (n + 2), cudaMemcpyHostToDevice);
207  cudaMemcpy(d_E_prev, E_prev, sizeof(double) * (m + 2) * (n + 2), cudaMemcpyHostToDevice);
208  cudaMemcpy(d_R, R, sizeof(double) * (m + 2) * (n + 2), cudaMemcpyHostToDevice);
```

**Figure 4**

At last, we copy Device arrays to Host and then free the CUDA memory.

**Version 2:**

- In version 2, we fused PDE and ODE loops into one loop in the kernel: see Figure 5

```

96 __global__ void singleKernel(double *E, double *E_prev, double *R, const int n, const int m, const double kk, const double dt, const double a, const double epsilon,
97 const double M1, const double M2, const double b, const double alpha) {
98     int i = threadIdx.x + 1;
99     int j = blockIdx.x + 1;
100     int index = j * (n + 2) + i;
101     E[index] = E_prev[index] + alpha * (E_prev[index + 1] + E_prev[index - 1] - 4 * E_prev[index] + E_prev[index + m + 2] + E_prev[index - (m + 2)]);
102     E[index] = E[index] - dt * (kk * E[index] * (E[index] - a) * (E[index] - 1) + E[index] * R[index]);
103     R[index] = R[index] + dt * (epsilon + M1 * R[index] / (E[index] + M2)) * (-R[index] - kk * E[index] * (E[index] - b - 1));
104 }
105
106 void simulate(double *E, double *E_prev, double *R,
107 const double alpha, const int n, const int m, const double kk,
108 const double dt, const double a, const double epsilon,
109 const double M1, const double M2, const double b) {
110
111     ghostKernel<<<1, n>>>(E_prev, n, m);
112     singleKernel<<<m, n>>>(E, E_prev, R, n, m, kk, dt, a, epsilon, M1, M2, b, alpha);
113 }

```

**Figure 5****Version 3:**

-In version 3, we changed our singleKernel so that the threads would access their corresponding place in array (calculated by thread index), so the threads don't keep referencing the arrays for calculating new values: see Figure 6

```

96 __global__ void singleKernel(double *E, double *E_prev, double *R,
97 const int n, const int m, const double kk,
98 const double dt, const double a, const double epsilon,
99 const double M1, const double M2, const double b, const double alpha) {
100     int i = threadIdx.x + 1;
101     int j = blockIdx.x + 1;
102     int index = j * (n + 2) + i;
103     double E_temp = E[index];
104     double R_temp = R[index];
105     E_temp = E_prev[index] + alpha * (E_prev[index + 1] + E_prev[index - 1] - 4 * E_prev[index] + E_prev[index + m + 2] + E_prev[index - (m + 2)]);
106     E_temp = E_temp - dt * (kk * E_temp * (E_temp - a) * (E_temp - 1) + E_temp * R_temp);
107     R_temp = R_temp + dt * (epsilon + M1 * R_temp / (E_temp + M2)) * (-R_temp - kk * E_temp * (E_temp - b - 1));
108     E[index] = E_temp;
109     R[index] = R_temp;
110 }

```

**Figure 6**

**Version 4:**

- We created a shared 2D memory, like stencil, to collect data from device accordingly to the threadIdx, then we used the shared variable for obtaining the host arrays. We used synchronized CUDA function to prevent the race condition. Then we call the function in the simulate with numBlock and blockSize, see Figure 7

```

96 __global__ void singleKernel(double *E, double *E_prev, double *R, const int n, const int m, const double kk,
97 const double dt, const double a, const double epsilon,
98 const double M1, const double M2, const double b, const double alpha) {
99
100 int x_thread = threadIdx.x, y_thread = threadIdx.y, x_block = blockIdx.x, y_block = blockIdx.y, x_blockDim = blockDim.x, y_blockDim = blockDim.y;
101 const int block_size = 16;
102
103 __shared__ double device_memory_array[block_size + 2][block_size + 2];
104
105 if(x_thread == 0) {
106     int index = (y_block * y_blockDim * (n + 2)) + (x_block * x_blockDim) + ((y_thread + 1) * (n + 2));
107     for (int j = 0; j < x_blockDim + 2; j++) {
108         device_memory_array[y_thread + 1][j] = E_prev[index + j];
109     }
110     if(y_thread == 0) {
111         int index = (y_block * y_blockDim * (n + 2)) + (x_block * x_blockDim);
112         for (int j = 0; j < x_blockDim + 2; j++) {
113             device_memory_array[0][j] = E_prev[index + j];
114         }
115     }
116     if(y_thread == 1) {
117         int index = (y_block * y_blockDim * (n + 2)) + (x_block * x_blockDim) + ((y_blockDim + 1) * (n + 2));
118         for (int j = 0; j < x_blockDim + 2; j++) {
119             device_memory_array[y_blockDim + 1][j] = E_prev[index + j];
120         }
121     }
122 }
123
124 int index = (y_block * y_blockDim * (n + 2)) + (x_block * x_blockDim) + (n + 2) + 1 + (y_thread * (n + 2) + x_thread);
125
126 __syncthreads();
127 double E_temp = E[index];
128 double R_temp = R[index];
129
130 E_temp = device_memory_array[y_thread + 1][x_thread + 1] + alpha * [
131     device_memory_array[y_thread + 1][x_thread + 2] +
132     device_memory_array[y_thread + 1][x_thread] - 4 * device_memory_array[y_thread + 1][x_thread + 1] + device_memory_array[y_thread + 2][x_thread + 1] +
133     device_memory_array[y_thread][x_thread + 1]);
134 E_temp = E_temp - dt * (kk * E_temp * (E_temp - a) * (E_temp - 1) + E_temp * R_temp);
135 R_temp = R_temp + dt * (epsilon + M1 * R_temp / (E_temp + M2)) * (-R_temp - kk * E_temp * (E_temp - b - 1));
136
137 __syncthreads();
138 E[index] = E_temp;
139 R[index] = R_temp;
140 }
141
142 void simulate(double *E, double *E_prev, double *R,
143 const double alpha, const int n, const int m, const double kk,
144 const double dt, const double a, const double epsilon,
145 const double M1, const double M2, const double b) {
146
147     const dim3 block_size(16,16);
148     const dim3 num_blocks(n / block_size.x, n / block_size.y);
149     ghostKernel<<<1, n>>>(E_prev, n, m);
150     singleKernel<<<num_blocks, block_size>>>(E, E_prev, R, n, m, kk, dt, a, epsilon, M1, M2, b, alpha);
151 }

```

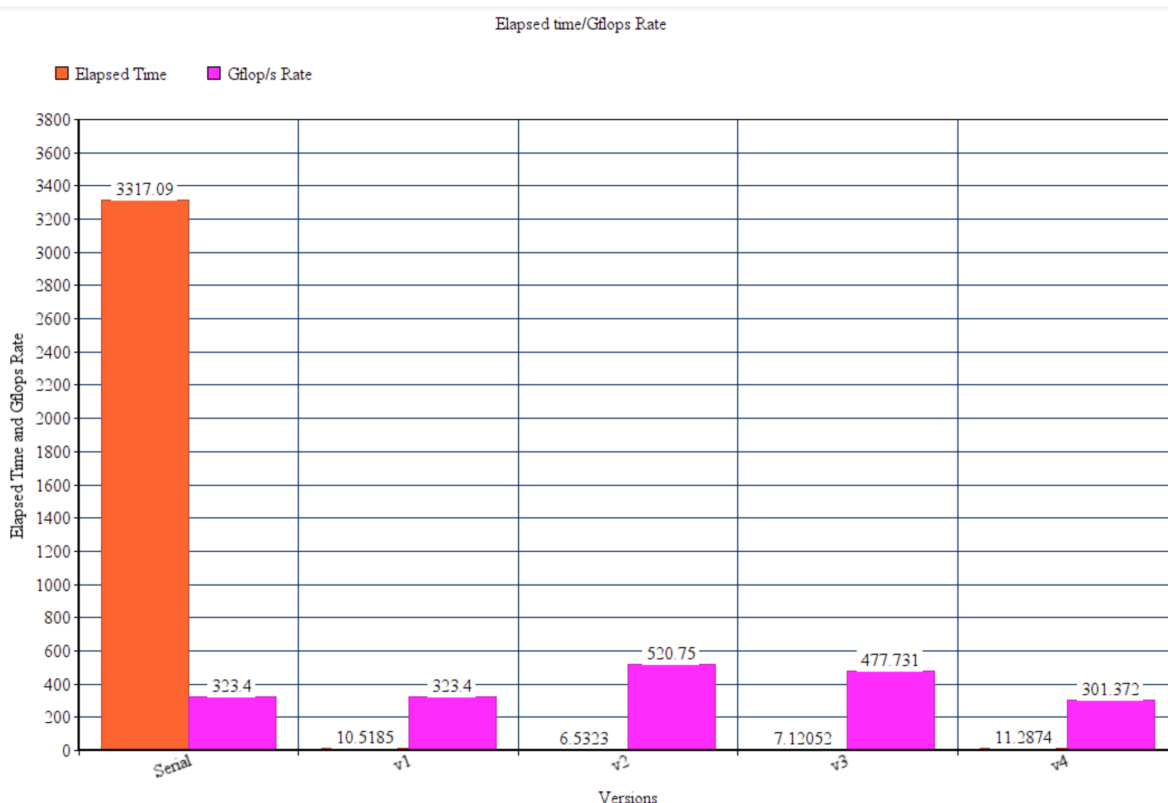
Figure 7

## Experiments

	Serial	Version 1	Version 2	Version 3	Version 4
Elapsed Time	3317.09	10.5185	6.5323	7.12052	11.2874
GFlop/s Rate	323.4	323.4	520.75	477.731	301.372

### Elapsed Time and Gflops Rate comparison:

-We measured Time and Gflop/s rate of our versions without using the plotter.  
Here are our result: **(Version 4 Block\_Size default value is 16)**



## Elapsed Time and Gflops Rate graph

- We have obtained close results for v1, v2, v3, and v4, but serial code was a lot more slower than our versions. Among the versions that have we created, v2 was always faster in every experiment that we have run. we have observed that, in the same version of the program: global memory > temp memory > shared memory in speed wise.

**Bandwidth Rate Comparison Experiment:**

-We used TESLA V100 GPU for all of our tests.

**v100 benchmark results:**

```
[CUDA Bandwidth Test] - Starting...
Running on...
```

```
Device 0: Tesla V100-SXM2-32GB
Quick Mode
```

```
Host to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                  12108.8
```

```
Device to Host Bandwidth, 1 Device(s)
PINNED Memory Transfers
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                  12862.3
```

```
Device to Device Bandwidth, 1 Device(s)
PINNED Memory Transfers
Transfer Size (Bytes)      Bandwidth(MB/s)
33554432                  731888.6
```



### Our version Bandwidths:

```
serial:  
Sustained Bandwidth (GB/sec): 1.17201
```

```
version 1:  
Sustained Bandwidth (GB/sec): 369.6
```

```
version 2:  
Sustained Bandwidth (GB/sec): 595.143
```

```
version 3:  
Sustained Bandwidth (GB/sec): 545.978
```

```
version 4:  
Sustained Bandwidth (GB/sec): 344.425
```

- Since v2 is the fastest, we have seen that its Sustained Bandwidth is also very high. As we have anticipated, even the v2 couldn't get too close to benchmark, since additional optimization is required.

### v4 Block Size Comparison:

#### Block size = 2

```
version 4(2x2):  
Elapsed Time (sec)      : 45.4588  
Sustained Gflops Rate   : 74.8302
```

#### Block size = 4

```
version 4(4x4):  
Elapsed Time (sec)      : 18.5759  
Sustained Gflops Rate   : 183.124
```

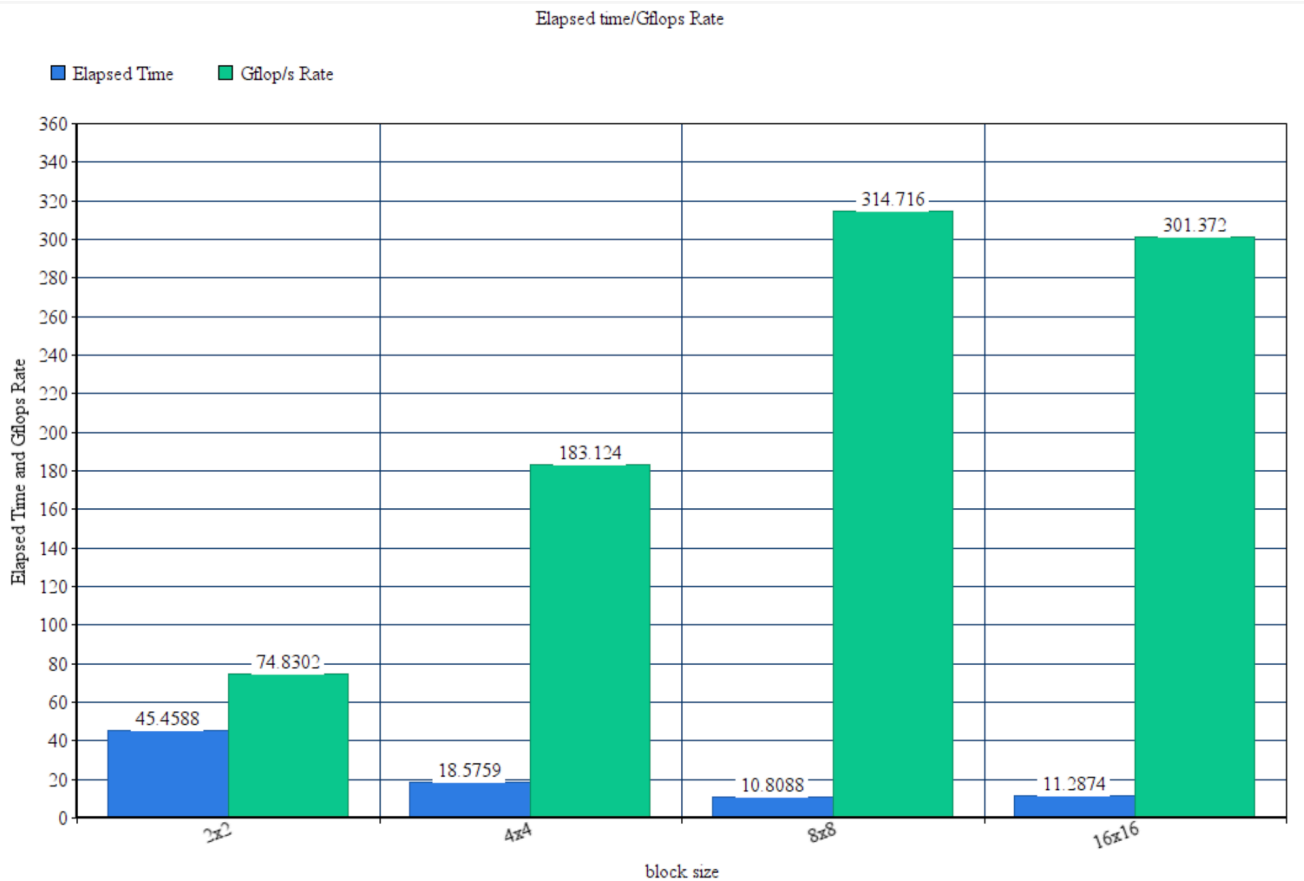
#### Block size = 8

```
version 4(8x8):  
Elapsed Time (sec)      : 10.8088  
Sustained Gflops Rate   : 314.716
```

#### Block size = 16

```
version 4(16x16):
```

```
Elapsed Time (sec)      : 11.2874
Sustained Gflops Rate   : 301.372
```



Elapsed Time and Gflops Rate vs Block size graph

- We observed that the best value is 8 for block size and 16 is very close to it performance wise. As we increase block size, from 2 to 8, we see performance increase because of overall less copy operation, but after 8, overhead becomes too much and performance starts to decrease again.